

RegressionAnalysis

April 1, 2023

Regression Analysis for the California Housing census

Author: Jose Pena

Github: JoseJuan98

0.1 Aims

The aim of this notebook is to demonstrate the data regression skills by applying them to a real-world dataset. A variety of tools will be leveraged to build and compare at least three linear regression models covering different variations, such as using a simple linear regression as a baseline, adding polynomial effects, and using a regularization regression. The focus should not only be on accuracy but also on explainability, as the models will be presented to a senior audience. In addition to presenting findings, recommendations for next steps in analyzing the data will be provided.

0.2 Instructions:

The purpose of this exercise is to demonstrate the ability to apply data regression skills using a variety of tools. The focus of this report should be on presenting findings, insights, and next steps. Visuals from code output may be included, but the emphasis should be on summarizing the findings rather than on reviewing the code. The expectation is that a wide range of tools will be leveraged to produce accurate and understandable results.

This exercise is an opportunity to demonstrate data regression skills. Follow the steps below to complete the analysis:

- A data set that is interesting and relevant to your work or field of study will be chosen. A brief description of the dataset, including its source, size, and the variables it contains will be provided.
- The data set will be explored and necessary cleaning and feature engineering will be performed. This will include, but is not limited to, checking for missing values, outliers, and correlations among variables. Visualizations or statistical tests may be used to support findings.
- At least three linear regression models that cover different variations, such as using a simple linear regression as a baseline, adding polynomial effects, and using a regularization regression, will be trained. The same training and test splits or the same cross-validation method will be used to ensure comparability.

- The performance of each model will be evaluated, and a recommendation for the best one based on accuracy and explainability will be provided. Metrics such as R-squared, mean squared error, or root mean squared error may be used to compare the models.
- The key findings and insights derived from the linear regression models will be summarized. The main drivers of the models and any insights gained from the data will be explained. Visualizations or tables will be used to support conclusions.
- Suggestions for next steps in analyzing the data will be provided. This could include exploring other variables or adding more data to improve the accuracy or explainability of the models.
- A report summarizing the analysis and findings will be prepared. The report will be geared towards a senior audience, such as a Chief Data Officer or Head of Analytics. Visualizations and tables will be included to support findings, but the focus will be on presenting insights and recommendations in a clear and concise manner.
- The report will be submitted for review by one of your peers. Feedback on the analysis will be received, as well as suggestions for improvement. This feedback will be used to revise the report as needed.

Remember to document the code and provide clear explanations of the methodology and thought process throughout the notebook. Good luck!

1 Setup

For this Regression Data Analysis project the following libraries will be used:

- `pandas` for managing the data.
- `numpy` for mathematical operations.
- `seaborn` for visualizing the data.
- `matplotlib` for visualizing the data.
- `sklearn` for machine learning related functions.

1.1 Imports

```
[1]: import pandas
import numpy

from matplotlib import pyplot
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, ElasticNetCV
from sklearn.pipeline import Pipeline
from sklearn.model_selection import GridSearchCV
from sklearn.preprocessing import PolynomialFeatures

from scipy.stats import normaltest
from scipy.special import inv_boxcox

from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```

# To be able to import outside of this folder
import os
os.chdir("..")

from src.utils.data_transformations import preprocess_data
from src.utils.evaluation import evaluate_metrics

# utils.py
from src.utils import fetch_housing_data, HOUSING_URL, HOUSING_PATH, TARGET

```

Setting up some options:

```

[2]: # Display and store plot within the notebook
    %matplotlib inline

    # Show all columns when displaying dataframe
    pandas.options.display.max_columns = None

```

Load data, if it wasn't downloaded before it will fetch it from the URL specified, otherwise will load it from a local 'csv' file.

```

[3]: data = fetch_housing_data(url=HOUSING_URL,
                               path=HOUSING_PATH)

```

2 1. About the Data

This California Housing dataset is available from [Luís Torgo's page](#) (University of Porto).

This dataset appeared in a 1997 paper titled Sparse Spatial Autoregressions by Pace, R. Kelley and Ronald Barry, published in the Statistics and Probability Letters journal. They built it using the 1990 California census data. It contains one row per census block group. A block group is the smallest geographical unit for which the U.S. Census Bureau publishes sample data (a block group typically has a population of 600 to 3,000 people).

The target variable or dependent variable for this analysis will be the `median_house_value`, which describes median price of the houses per block group.

Shape of the dataset

```

[4]: data.shape

```

```

[4]: (20640, 10)

```

List of columns

```

[5]: data.columns.to_list()

```

```

[5]: ['longitude',
      'latitude',

```

```
'housing_median_age',
'total_rooms',
'total_bedrooms',
'population',
'households',
'median_income',
'median_house_value',
'ocean_proximity']
```

First 5 rows of the dataset

```
[6]: data.head()
```

```
[6]:   longitude  latitude  housing_median_age  total_rooms  total_bedrooms  \
0    -122.23    37.88             41.0         880.0         129.0
1    -122.22    37.86             21.0        7099.0        1106.0
2    -122.24    37.85             52.0        1467.0         190.0
3    -122.25    37.85             52.0        1274.0         235.0
4    -122.25    37.85             52.0        1627.0         280.0

      population  households  median_income  median_house_value  ocean_proximity
0         322.0        126.0         8.3252         452600.0        NEAR BAY
1        2401.0       1138.0         8.3014        358500.0        NEAR BAY
2         496.0        177.0         7.2574        352100.0        NEAR BAY
3         558.0        219.0         5.6431        341300.0        NEAR BAY
4         565.0        259.0         3.8462        342200.0        NEAR BAY
```

Non-null values count, type of feature and memory usage

```
[7]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude             20640 non-null  float64
1   latitude              20640 non-null  float64
2   housing_median_age    20640 non-null  float64
3   total_rooms           20640 non-null  float64
4   total_bedrooms        20433 non-null  float64
5   population            20640 non-null  float64
6   households            20640 non-null  float64
7   median_income         20640 non-null  float64
8   median_house_value    20640 non-null  float64
9   ocean_proximity       20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

Statistical properties of the dataset

```
[8]: data.describe(include='all')
```

```
[8]:
```

	longitude	latitude	housing_median_age	total_rooms	\
count	20640.000000	20640.000000	20640.000000	20640.000000	
unique	NaN	NaN	NaN	NaN	
top	NaN	NaN	NaN	NaN	
freq	NaN	NaN	NaN	NaN	
mean	-119.569704	35.631861	28.639486	2635.763081	
std	2.003532	2.135952	12.585558	2181.615252	
min	-124.350000	32.540000	1.000000	2.000000	
25%	-121.800000	33.930000	18.000000	1447.750000	
50%	-118.490000	34.260000	29.000000	2127.000000	
75%	-118.010000	37.710000	37.000000	3148.000000	
max	-114.310000	41.950000	52.000000	39320.000000	

	total_bedrooms	population	households	median_income	\
count	20433.000000	20640.000000	20640.000000	20640.000000	
unique	NaN	NaN	NaN	NaN	
top	NaN	NaN	NaN	NaN	
freq	NaN	NaN	NaN	NaN	
mean	537.870553	1425.476744	499.539680	3.870671	
std	421.385070	1132.462122	382.329753	1.899822	
min	1.000000	3.000000	1.000000	0.499900	
25%	296.000000	787.000000	280.000000	2.563400	
50%	435.000000	1166.000000	409.000000	3.534800	
75%	647.000000	1725.000000	605.000000	4.743250	
max	6445.000000	35682.000000	6082.000000	15.000100	

	median_house_value	ocean_proximity
count	20640.000000	20640
unique	NaN	5
top	NaN	<1H OCEAN
freq	NaN	9136
mean	206855.816909	NaN
std	115395.615874	NaN
min	14999.000000	NaN
25%	119600.000000	NaN
50%	179700.000000	NaN
75%	264725.000000	NaN
max	500001.000000	NaN

As it was analyzed in the previous exercise of Exploratory Data Analysis ([notebook](#), [report](#)) the actions taken for Data Cleaning and Feature Engineering are: * Target normalization * Handling missing values * Handling outliers * Encoding categorical variables * Scaling continuous variables

And these actions are encapsulated in the method `prepare_data()` from the `utils/data_transformations.py`, but first the data must be split to create the train and

test set to avoid overfitting and inaccurate evaluation.

```
[9]: x_train, x_test, y_train, y_test = train_test_split(data.drop(columns=[TARGET],  
↪axis=1), data[TARGET], random_state=42, test_size=0.3)

[10]: f" Shape x_train {x_train.shape} - y_train {y_train.shape}"

[10]: ' Shape x_train (14448, 9) - y_train (14448,)'

[11]: f" Shape x_test {x_test.shape} - y_test {y_test.shape}"

[11]: ' Shape x_test (6192, 9) - y_test (6192,)'
```

3 2. Objectives

This exercise focuses in the predictions of the models, so this approach compares y_p with y by **performance metrics**, which measure the quality of the model's predictions (closeness between y_p and y).

As this approach doesn't focus on interpretability there is a greater risk of having a Black-box model, so it's recommended also to explore an approach based in interpretation to have both.

4 3. Linear Regression Models

```
[12]: metrics = {}

def add_metrics(metrics_report: dict, model_name: str, y_true, y_pred, lambda:  
↪float) -> dict:
    """
    Args:
        metrics_report:
        model_name:
        y_true:
        y_pred:

    Returns:
        dict: per model
            - MSE: penalizes big errors
            - RMSE: standarize unit errors
            - R2: proportion of variance (0,1) - The bigger better
            - MAE: average of erros

    """
    metrics_report[model_name] = evaluate_metrics(y_true=y_true, y_pred=y_pred,  
↪lambda=lambda)
    return metrics_report
```

4.1 Preparing the data

4.1.1 Multicollinearity Analysis

As observed in the Exploratory Data Analysis phase, it's likely that there is multicollinearity between some variables.

Correlation Matrix (Pearson's) To check it, first it can be observed the correlation matrix of the numerical values by computing the Person's correlation.

```
[13]: # Compute correlation matrix
corr_matrix = x_train.corr(numeric_only=True, method="pearson")

print(corr_matrix)
```

	longitude	latitude	housing_median_age	total_rooms	\
longitude	1.000000	-0.923408	-0.101083	0.037158	
latitude	-0.923408	1.000000	0.003461	-0.028768	
housing_median_age	-0.101083	0.003461	1.000000	-0.362713	
total_rooms	0.037158	-0.028768	-0.362713	1.000000	
total_bedrooms	0.061797	-0.059700	-0.321328	0.929527	
population	0.092163	-0.101665	-0.291589	0.855384	
households	0.047659	-0.063487	-0.302516	0.920133	
median_income	-0.019019	-0.075892	-0.117506	0.198362	

	total_bedrooms	population	households	median_income
longitude	0.061797	0.092163	0.047659	-0.019019
latitude	-0.059700	-0.101665	-0.063487	-0.075892
housing_median_age	-0.321328	-0.291589	-0.302516	-0.117506
total_rooms	0.929527	0.855384	0.920133	0.198362
total_bedrooms	1.000000	0.876119	0.980570	-0.010193
population	0.876119	1.000000	0.904678	0.003661
households	0.980570	0.904678	1.000000	0.011628
median_income	-0.010193	0.003661	0.011628	1.000000

As shown, there are clearly some correlations within the dependent variables which can cause collinearity.

Variance Inflation Factor To check if they really cause multicollinearity the Variance Inflation Factor can be used. Because of the scale used in the following method, If it gives a value bigger than 15 or a relatively bigger value compare with the rest it can be proved that those variables will definitely cause collinearity.

```
[14]: numerical_X_train = x_train.select_dtypes(include='number')

# Compute VIF scores
vif_scores = pandas.Series([variance_inflation_factor(numerical_X_train.values,
↪ i) for i in range(numerical_X_train.shape[1])], index=numerical_X_train.
↪ columns)
```

```
# Print VIF scores
print(vif_scores)
```

```
longitude      614.104270
latitude       548.596876
housing_median_age  7.259158
total_rooms     30.774710
total_bedrooms  95.973312
population     15.813275
households     94.392021
median_income   8.260040
dtype: float64
```

```
[15]: corr_matrix[corr_matrix > 0.7]
```

```
[15]:
```

	longitude	latitude	housing_median_age	total_rooms	\
longitude	1.0	NaN	NaN	NaN	
latitude	NaN	1.0	NaN	NaN	
housing_median_age	NaN	NaN	1.0	NaN	
total_rooms	NaN	NaN	NaN	1.000000	
total_bedrooms	NaN	NaN	NaN	0.929527	
population	NaN	NaN	NaN	0.855384	
households	NaN	NaN	NaN	0.920133	
median_income	NaN	NaN	NaN	NaN	

	total_bedrooms	population	households	median_income
longitude	NaN	NaN	NaN	NaN
latitude	NaN	NaN	NaN	NaN
housing_median_age	NaN	NaN	NaN	NaN
total_rooms	0.929527	0.855384	0.920133	NaN
total_bedrooms	1.000000	0.876119	0.980570	NaN
population	0.876119	1.000000	0.904678	NaN
households	0.980570	0.904678	1.000000	NaN
median_income	NaN	NaN	NaN	1.0

Based on the results from the VIF score and correlation matrix, it can be observed that **total_rooms** and **total_bedrooms** have high values of VIF, indicating that these two variables are highly correlated. This high correlation is not surprising, as **total_rooms** may affect the number of **total_bedrooms**. However, including both features in the model can lead to issues with multicollinearity, as it becomes difficult to distinguish the individual effect of each variable on the target variable. Therefore, it may be necessary to address multicollinearity in the model, either by removing one of the highly correlated features or by using techniques such as ridge regression or principal component analysis.

The same with **population** and **households**.


```
[16]: vars_to_remove = ['total_rooms', 'population', 'households']
x_train = x_train.drop(columns=vars_to_remove, axis=1)
x_test = x_test.drop(columns=vars_to_remove, axis=1)

[17]: vars_with_outliers = ["median_income", "total_bedrooms",
                           "housing_median_age"]

x_train, y_train, preprocessor, lambda = preprocess_data(X=x_train, y=y_train,
↳variables_with_outliers=vars_with_outliers, normalize_target=True)
x_test, y_test, _, _ = preprocess_data(X=x_test, y=y_test,
↳preprocessor=preprocessor, variables_with_outliers=vars_with_outliers)

[18]: preprocessor

[18]: ColumnTransformer(n_jobs=-1, remainder='passthrough',
                        transformers=[('numerical',
                                      Pipeline(steps=[('num_imputer',
SimpleImputer(missing_values=NA)),
                                                    ('scaler', MinMaxScaler())],
                                      verbose=True),
                                      ['longitude', 'latitude', 'housing_median_age',
                                      'total_bedrooms', 'median_income']),
                                      ('categorical',
                                      Pipeline(steps=[('cat_imputer',
SimpleImputer(strategy='most_frequent')),
                                                    ('cat_ohe',
OneHotEncoder(handle_unknown='ignore'))],
                                      verbose=True),
                                      ['ocean_proximity'])]],
                        verbose=True)
```

4.1.2 Normality Test in the target variable

The `normaltest()` function is a statistical test for normality that combines skewness and kurtosis based on D'Agostino and Pearson's method.

It produces a p-value, which indicates the goodness of fit to a normal distribution. A higher p-value suggests a closer match to a normal distribution. Generally, frequentist statisticians consider a p-value greater than 0.05 as evidence that the distribution is normal, and fail to reject the null hypothesis of normality.

However, it's important to note that this test is not perfect and has some limitations.

```
[19]: normaltest(y_train)

[19]: NormaltestResult(statistic=array([332.66882662]),
pvalue=array([5.77939196e-73]))

[20]: normaltest(y_test)
```

```
[20]: NormaltestResult(statistic=716.5809435890789, pvalue=2.4912951418262196e-156)
```

4.1.3 Simple Linear Regression

```
[21]: lr = GridSearchCV(estimator=LinearRegression(n_jobs=-1),
                        n_jobs=-1,
                        verbose=1,
                        param_grid={},
                        cv=5
                        )
```

```
[22]: lr.fit(X=x_train, y=y_train)
```

Fitting 5 folds for each of 1 candidates, totalling 5 fits

```
[22]: GridSearchCV(cv=5, estimator=LinearRegression(n_jobs=-1), n_jobs=-1,
                  param_grid={}, verbose=1)
```

```
[23]: metrics = add_metrics(metrics_report=metrics,
                             model_name='LR_Simple',
                             y_true=y_test,
                             y_pred=lr.predict(x_test),
                             lambda=lambda)
```

```
print(f"{metrics['LR_Simple']}")
```

```
{'MSE': 84154.5075, 'RMSE': 290.094, 'R2': 0.4604, 'MAE': 56842.5729}
```

4.1.4 Linear Regression with Polynomial Features

```
[24]: model_polyn_pipe = Pipeline(steps=[
      ('polynomail', PolynomialFeatures(degree=3, include_bias=False)),
      ('lr', LinearRegression(n_jobs=-1))
    ])
```

```
lr_polynomial = GridSearchCV(
    estimator=model_polyn_pipe,
    n_jobs=-1,
    verbose=1,
    param_grid={},
    cv=5
)
```

```
[25]: lr_polynomial.fit(X=x_train, y=y_train)
```

Fitting 5 folds for each of 1 candidates, totalling 5 fits

```
[25]: GridSearchCV(cv=5,
                  estimator=Pipeline(steps=[('polynomail',
```

```

        PolynomialFeatures(degree=3,
                           include_bias=False)),
        ('lr', LinearRegression(n_jobs=-1))]),
    n_jobs=-1, param_grid={}, verbose=1)

```

```

[26]: metrics = add_metrics(metrics_report=metrics,
                            model_name='LR_PolynEffects',
                            y_true=y_test,
                            y_pred=lr_polynomial.predict(x_test),
                            lambda=lambda)

print(f"{metrics['LR_PolynEffects']}")

```

```
{'MSE': 76401.5438, 'RMSE': 276.4083, 'R2': 0.555, 'MAE': 50224.9909}
```

4.1.5 Regression with Regularization

```

[27]: l1_ratios = numpy.linspace(0.1, 0.9, 9)
      alphas = numpy.array([1e-5, 5e-5, 0.0001, 0.0005, 0.1, 0.01, 1])

      lr_regularization = ElasticNetCV(
          alphas=alphas,
          l1_ratio=l1_ratios,
          n_jobs=-1,
          max_iter=int(1e5),
          cv=5
      )

```

```

[28]: lr_regularization.fit(X=x_train, y=y_train.to_numpy().ravel())

```

```

[28]: ElasticNetCV(alphas=array([1.e-05, 5.e-05, 1.e-04, 5.e-04, 1.e-01, 1.e-02,
1.e+00])),
          cv=5,
          l1_ratio=array([0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9]),
          max_iter=100000, n_jobs=-1)

```

```

[29]: metrics = add_metrics(metrics_report=metrics,
                            model_name='Regularization',
                            y_true=y_test,
                            y_pred=lr_regularization.predict(x_test),
                            lambda=lambda)

print(f"{metrics['Regularization']}")

```

```
{'MSE': 84157.8672, 'RMSE': 290.0998, 'R2': 0.4604, 'MAE': 56845.0172}
```

```

[30]: import json
      print(json.dumps(metrics, indent=4))

```

```
{
  "LR_Simple": {
    "MSE": 84154.5075,
    "RMSE": 290.094,
    "R2": 0.4604,
    "MAE": 56842.5729
  },
  "LR_PolynEffects": {
    "MSE": 76401.5438,
    "RMSE": 276.4083,
    "R2": 0.555,
    "MAE": 50224.9909
  },
  "Regularization": {
    "MSE": 84157.8672,
    "RMSE": 290.0998,
    "R2": 0.4604,
    "MAE": 56845.0172
  }
}
```

The Linear Regression with Polynomial Effects seems to be the best fit as it has the lowest MSE and RMSE, highest R2, and lowest MAE. Additionally, using polynomial features can help capture more complex relationships between the features and target variable. However, it's always good to consider the interpretability of the model, and the simple linear regression model could be preferred if interpretability is a priority.

5 4. Insights and key findings

Given the properties of the test set:

```
[31]: print(f'Mean: {y_test.mean()}\nMedian: {y_test.median()}\nMinimum: {y_test.\n      ↪min()}\nMaximum: {y_test.max()}')
```

```
Mean: 206696.8142764858
Median: 181000.0
Minimum: 14999.0
Maximum: 500001.0
```

Regarding the key findings and insights from the linear regression model, we can observe the following:

- The models explain around 46%-55% of the variance in the target variable, which is moderate to good.
- The RMSE values suggest that the models have an average error of around \$276K-\$290K in predicting the target variable.
- The MAE values suggest that, on average, the predicted values are off by around \$50K-\$56K.

- The mean and median values of the target variable suggest that the dataset has a right-skewed distribution.
- The minimum and maximum values of the target variable indicate that there are significant differences between the lowest and highest values.

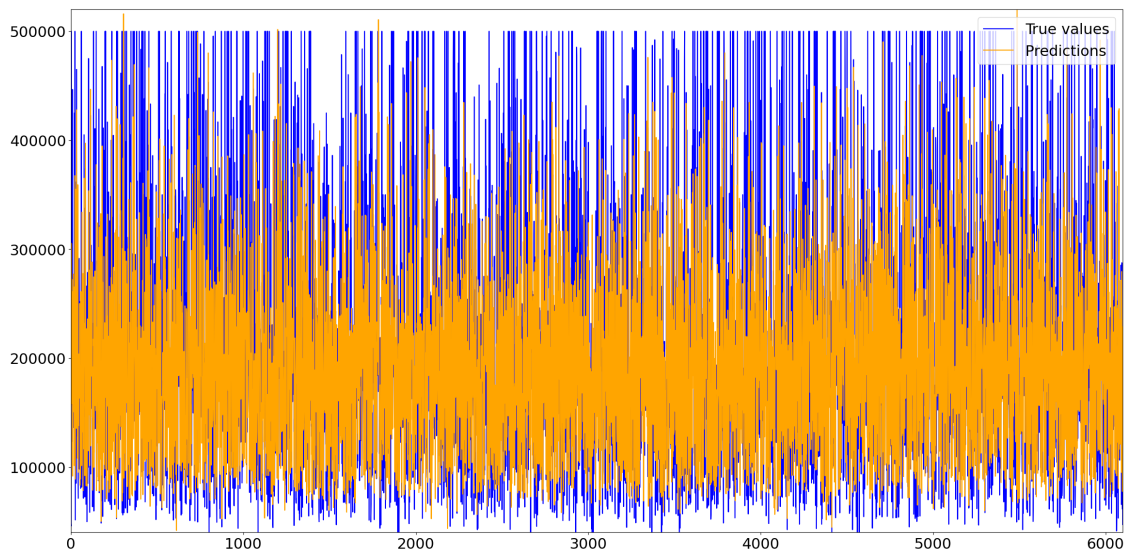
It's essential to note that the models assume linearity between the features and target variable. Still, this assumption may not hold in some cases, and nonlinear models could perform better.

5.1 Insights Visualization

5.1.1 Actual vs Predicted Values visualization

```
[32]: # Inverse the normalization to get easier interpretations of the results
y_pred = inv_boxcox(lr_regularization.predict(x_test), lambda)

pyplot.figure(figsize=(28,14))
pyplot.plot(range(len(y_test)), y_test, color='blue', label='True values')
pyplot.plot(range(len(y_pred)), y_pred, color='orange', label='Predictions')
pyplot.xticks(fontsize=22)
pyplot.yticks(fontsize=22)
pyplot.legend(fontsize=22)
pyplot.ylim(40000, 520000)
pyplot.xlim(0, 6100)
pyplot.show();
```

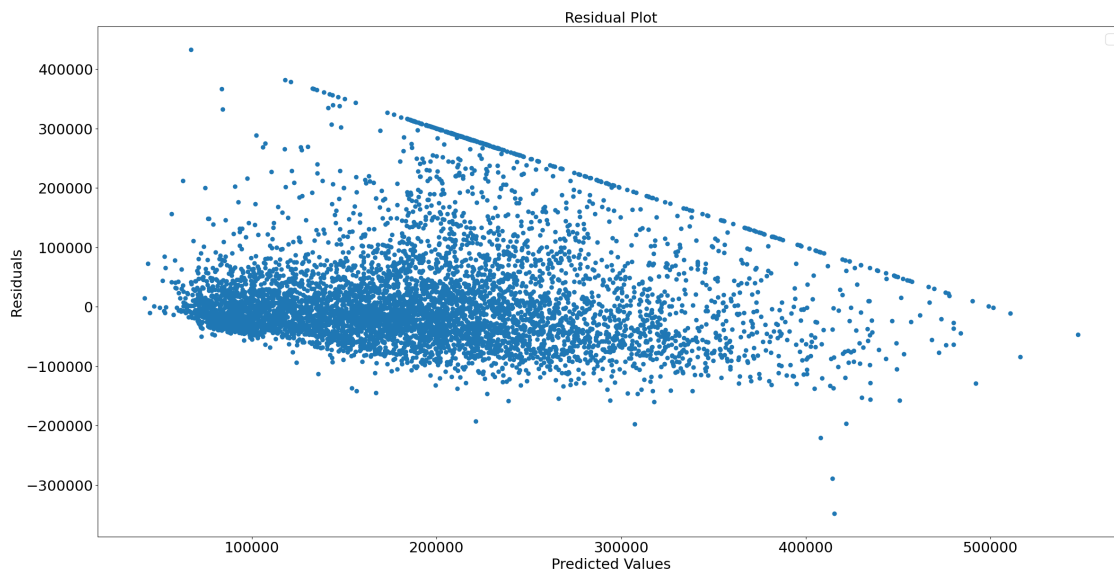


5.1.2 Residual Plot

```
[33]: # Compute the residuals
residuals = numpy.subtract(y_test, y_pred)

# Plot the residuals
pyplot.figure(figsize=(28,14))
pyplot.scatter(y_pred, residuals)
pyplot.title('Residual Plot', fontsize=22)
pyplot.xlabel('Predicted Values', fontsize=22)
pyplot.ylabel('Residuals', fontsize=22)
pyplot.xticks(fontsize=22)
pyplot.yticks(fontsize=22)
pyplot.legend(fontsize=22)
pyplot.show();
```

No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.



5.1.3 Feature Importances Plot

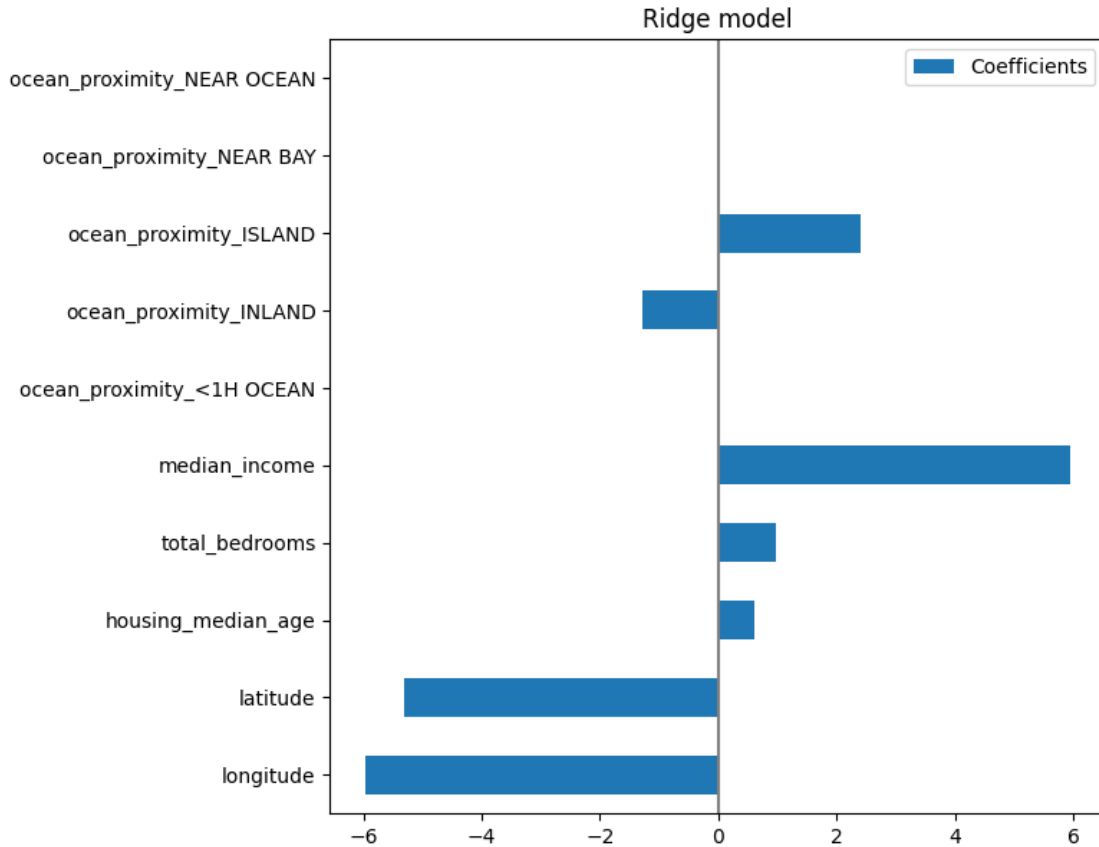
Creating a bar chart of the feature importances by plotting the coefficients of the linear regression model can help to identify which features are most important in predicting the response variable.

```
[34]: coefs = pandas.DataFrame(
    lr_regularization.coef_,
    columns=['Coefficients'], index=x_train.columns
)
```

```

coefs.plot(kind='barh', figsize=(9, 7))
pyplot.title('Ridge model')
pyplot.axvline(x=0, color='.5')
pyplot.subplots_adjust(left=.3)

```



6 5. Next Steps

Regarding next steps, we could explore several options:

- Feature engineering: We could try to identify additional relevant features that may explain more variance in the target variable and improve the model's performance.
- Nonlinear models: As mentioned earlier, we could explore more sophisticated models that can capture nonlinear relationships between the features and target variable.
- Outlier detection: We could analyze further if there are any outliers in the dataset that could affect the model's performance and try to remove or transform them if necessary.
- Cross-validation: We could apply cross-validation to evaluate the models' performance and ensure that they generalize well to new data.
- Ensemble methods: We could consider ensemble methods, such as random forests or gradient boosting, that can combine multiple models and improve their predictive power.

Overall, selecting the best model depends on the specific goals and priorities of the analysis, and additional exploration of the data and modeling techniques could provide further insights and improvements.

6.1 Author

Author: Jose Pena

Github: JoseJuan98

6.2 Change Log

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2023-01-27	1.0	Jose	First version